

# Instructions

- Some parts of the code are already done for you
- You need to execute all the cells
- You need to add the code where ever you see "#### Add your code here ####"
- Marks are mentioned along with the cells

## ▼ Face detection

Task is to predict the boundaries(mask) around the face in a given image.

## ▼ Dataset

Faces in images marked with bounding boxes. Have around 500 images with around 1100 faces mar

## ▼ Mount Google drive if you are using google colab

- We recommend using Google Colab as you can face memory issues and longer runtimes while

```
from google.colab import drive
drive.mount('/content/drive')
```

🔗 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mour

## ▼ Change current working directory to project folder (1 mark)

Saved successfully! ✕

```
import os
import tensorflow as tf
#### Add your code here ####
project_dir = "/content/drive/My Drive/greatlakes/Projects/Advanced_Computer_Vision/Project1/"
os.chdir(project_dir)
```

## ▼ Load the "images.npy" file (2 marks)

- This file contains images with details of bounding boxes

```
import numpy as np
data = np.load('images.npy', allow_pickle=True)
```

## ▼ Check one sample from the loaded "images.npy" file (2 marks)

Hint - print data[10][1]

```
#### Add your code here ####
```

```
print(data[10][1])
```

```
↳ [{'label': ['Face'], 'notes': '', 'points': [{'x': 0.48, 'y': 0.10385756676557864}, {'x':
```

## ▼ Set image dimensions (1 mark)

- Initialize image height, image width with value: 224

```
IMAGE_WIDTH = 224
```

```
IMAGE_HEIGHT = 224
```

## ▼ Create features and labels

- Here feature is the image
- The label is the mask
- Images will be stored in "X\_train" array
- Masks will be stored in "masks" array

```
import cv2
```

```
from tensorflow.keras.applications.mobilenet import preprocess_input
```

```
masks = np.zeros((int(data.shape[0]), IMAGE_HEIGHT, IMAGE_WIDTH))
```

```
for index in range(int(data.shape[0])):
    masks[index] = np.zeros((IMAGE_HEIGHT, IMAGE_WIDTH, 3))
```

Saved successfully!

```
img = cv2.resize(img, dsize=(IMAGE_HEIGHT, IMAGE_WIDTH), interpolation=cv2.INTER_CUBIC)
```

```
try:
```

```
    img = img[:, :, :3]
```

```
except:
```

```
    continue
```

```
X_train[index] = preprocess_input(np.array(img, dtype=np.float32))
```

```
for i in data[index][1]:
```

```
    x1 = int(i["points"][0]['x'] * IMAGE_WIDTH)
```

```
    x2 = int(i["points"][1]['x'] * IMAGE_WIDTH)
```

```
    y1 = int(i["points"][0]['y'] * IMAGE_HEIGHT)
```

```
    y2 = int(i["points"][1]['y'] * IMAGE_HEIGHT)
```

```
    masks[index][y1:y2, x1:x2] = 1
```

### ▼ Print the shape of X\_train and mask array (1 mark)

```
X_train.shape
```

```
↳ (409, 224, 224, 3)
```

```
masks.shape
```

```
↳ (409, 224, 224)
```

### ▼ Print a sample image and image array

```
from matplotlib import pyplot  
n = 10  
print(X_train[n])  
pyplot.imshow(X_train[n])
```

```
↳
```

Saved successfully!



```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0
[[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 ...
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]]

[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 ...
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]]

[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 ...
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]]

...

[[-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 ...
 [-0.96862745 -0.96862745 -0.96862745]
 [-0.96078432 -0.96078432 -0.96078432]
 [-0.96078432 -0.96078432 -0.96078432]]


[[-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]


...
[-0.96862745 -0.96862745 -0.96862745]
[-0.96078432 -0.96078432 -0.96078432]
[-0.95294118 -0.95294118 -0.95294118]]

[[-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 [-1.          -1.          -1.          ]
 ...
 [-0.97647059 -0.97647059 -0.97647059]
 [-0.96862745 -0.96862745 -0.96862745]
 [-0.96078432 -0.96078432 -0.96078432]]]
<matplotlib.image.AxesImage at 0x7fcea46a98d0>

```


Saved successfully!





```
pyplot.imshow(masks[n])
```

```
<matplotlib.image.AxesImage at 0x7fcea4739390>
```



## ▼ Create the model (10 marks)

- Add MobileNet as model with below parameter values
  - input\_shape: IMAGE\_HEIGHT, IMAGE\_WIDTH, 3
  - include\_top: False
  - alpha: 1.0

Saved successfully!

- Add UNET architecture layers
  - This is the trickiest part of the project, you need to research and implement it correctly

```
from tensorflow.keras.applications.mobilenet import MobileNet
from tensorflow.keras.layers import Concatenate, UpSampling2D, Conv2D, Reshape
from tensorflow.keras.models import Model
```

ALPHA = 1.0 # Width hyper parameter for MobileNet (0.25, 0.5, 0.75, 1.0). Higher width means

```
def create_model(trainable=True):
    model = MobileNet(input_shape=(IMAGE_WIDTH, IMAGE_HEIGHT, 3), include_top=False,
    for layer in model.layers:
        layer.trainable = trainable
```

```
# Add all the UNET layers here
```

```
block1 = model.get_layer("conv_pw_1_relu").output
block2 = model.get_layer("conv_pw_3_relu").output
block3 = model.get_layer("conv_pw_5_relu").output
block4 = model.get_layer("conv_pw_11_relu").output
block5 = model.get_layer("conv_pw_13_relu").output
```

```
x = Concatenate()([UpSampling2D()(block5), block4])
x = Concatenate()([UpSampling2D()(x), block3])
x = Concatenate()([UpSampling2D()(x), block2])
x = Concatenate()([UpSampling2D()(x), block1])
x = UpSampling2D()(x)
```

```
x = Conv2D(1, kernel_size=1, activation="sigmoid")(x)
x = Reshape((IMAGE_WIDTH, IMAGE_HEIGHT))(x)
```

```
return Model(inputs=model.input, outputs=x) ##### Add your code here #####
```

## ▼ Call the create\_model function

```
# Give trainable=False as argument, if you want to freeze lower layers for fast training (but
model = create_model()
```

```
# Print summary
model.summary()
```



Saved successfully!



conv_pw_10_bn (BatchNormalizati	(None, 14, 14, 512)	2048	conv_pw_10[0][0]
conv_pw_10_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_10_bn[0][0]
conv_dw_11 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_10_relu[0][0]
conv_dw_11_bn (BatchNormalizati	(None, 14, 14, 512)	2048	conv_dw_11[0][0]
conv_dw_11_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_11_bn[0][0]
conv_pw_11 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_11_relu[0][0]
conv_pw_11_bn (BatchNormalizati	(None, 14, 14, 512)	2048	conv_pw_11[0][0]
conv_pw_11_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_11_bn[0][0]
conv_pad_12 (ZeroPadding2D)	(None, 15, 15, 512)	0	conv_pw_11_relu[0][0]
conv_dw_12 (DepthwiseConv2D)	(None, 7, 7, 512)	4608	conv_pad_12[0][0]
conv_dw_12_bn (BatchNormalizati	(None, 7, 7, 512)	2048	conv_dw_12[0][0]
conv_dw_12_relu (ReLU)	(None, 7, 7, 512)	0	conv_dw_12_bn[0][0]
conv_pw_12 (Conv2D)	(None, 7, 7, 1024)	524288	conv_dw_12_relu[0][0]
conv_pw_12_bn (BatchNormalizati	(None, 7, 7, 1024)	4096	conv_pw_12[0][0]
conv_pw_12_relu (ReLU)	(None, 7, 7, 1024)	0	conv_pw_12_bn[0][0]
conv_dw_13 (DepthwiseConv2D)	(None, 7, 7, 1024)	9216	conv_pw_12_relu[0][0]
conv_dw_13_bn (BatchNormalizati	(None, 7, 7, 1024)	4096	conv_dw_13[0][0]
conv_dw_13_relu (ReLU)	(None, 7, 7, 1024)	0	conv_dw_13_bn[0][0]
conv_pw_13 (Conv2D)	(None, 7, 7, 1024)	1048576	conv_dw_13_relu[0][0]
conv_pw_13_relu (ReLU)	(None, 7, 7, 1024)	4096	conv_pw_13[0][0]
conv_pw_13_relu (ReLU)	(None, 7, 7, 1024)	0	conv_pw_13_bn[0][0]
up_sampling2d_5 (UpSampling2D)	(None, 14, 14, 1024)	0	conv_pw_13_relu[0][0]
concatenate_4 (Concatenate)	(None, 14, 14, 1536)	0	up_sampling2d_5[0][0] conv_pw_11_relu[0][0]
up_sampling2d_6 (UpSampling2D)	(None, 28, 28, 1536)	0	concatenate_4[0][0]
concatenate_5 (Concatenate)	(None, 28, 28, 1792)	0	up_sampling2d_6[0][0] conv_pw_5_relu[0][0]
up_sampling2d_7 (UpSampling2D)	(None, 56, 56, 1792)	0	concatenate_5[0][0]
concatenate_6 (Concatenate)	(None, 56, 56, 1920)	0	up_sampling2d_7[0][0] conv_pw_3_relu[0][0]

Saved successfully!



up_sampling2d_8 (UpSampling2D)	(None, 112, 112, 192 0	concatenate_6[0][0]
concatenate_7 (Concatenate)	(None, 112, 112, 198 0	up_sampling2d_8[0][0] conv_pw_1_relu[0][0]
up_sampling2d_9 (UpSampling2D)	(None, 224, 224, 198 0	concatenate_7[0][0]
conv2d_1 (Conv2D)	(None, 224, 224, 1) 1985	up_sampling2d_9[0][0]
reshape_1 (Reshape)	(None, 224, 224) 0	conv2d_1[0][0]

=====

Total params: 3,230,849  
 Trainable params: 3,208,961  
 Non-trainable params: 21,888

Saved successfully!





Saved successfully!



## Define dice coefficient function (5 marks)

- Create a function to calculate dice coefficient

Saved successfully!

### ▼ Dice Coefficient (F1 Score) Explanation

The Dice Coefficient is  $2 * \text{the Area of Overlap}$  divided by the total number of pixels in both images

```
def dice_coefficient(y_true, y_pred):  
    ##### Add your code here #####  
    numerator = 2 * tf.reduce_sum(y_true * y_pred)  
    denominator = tf.reduce_sum(y_true + y_pred)  
  
    return numerator / (denominator + tf.keras.backend.epsilon())
```

### ▼ Define loss

```

from tensorflow.keras.losses import binary_crossentropy
from tensorflow.keras.backend import log, epsilon
def loss(y_true, y_pred):
    return binary_crossentropy(y_true, y_pred) - log(dice_coefficient(y_true, y_pred) + epsilon)

```

## ▼ Compile the model (2 marks)

- Compile the model using below parameters
  - loss: use the loss function defined above
  - optimizers: use Adam optimizer
  - metrics: use dice\_coefficient function defined above

```

#### Add your code here ####
from tensorflow.keras.optimizers import Adam
optimizer = Adam(lr=1e-4, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)
model.compile(loss=loss, optimizer=optimizer, metrics=[dice_coefficient])

```

## ▼ Define checkpoint and earlystopping

```

from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLRonPlateau
checkpoint = ModelCheckpoint("model-{loss:.2f}.h5", monitor="loss", verbose=1, save_best_only=True,
                             save_weights_only=True, mode="min", period=1)
stop = EarlyStopping(monitor="loss", patience=5, mode="min")
reduce_lr = ReduceLRonPlateau(monitor="loss", factor=0.2, patience=5, min_lr=1e-6, verbose=1,

```

⚠ WARNING:tensorflow:`period` argument is deprecated. Please use `save\_freq` to specify the

## ▼ Fit the model (2 marks)

Saved successfully! ✕ Parameters

- epochs: you can decide
- batch\_size: 1
- callbacks: checkpoint, reduce\_lr, stop

```

#### Add your code here ####
EPOCHS = 10
BATCH_SIZE = 1
model.fit(X_train, masks, batch_size=BATCH_SIZE, nb_epoch=EPOCHS, callbacks=[checkpoint, reduce_lr, stop],
          use_multiprocessing=False)

```

⚠

WARNING:tensorflow:The `nb\_epoch` argument in `fit` has been renamed `epochs`.

Train on 409 samples

Epoch 1/10

408/409 [=====>.] - ETA: 0s - loss: 1.2902 - dice\_coefficient: 0.

Epoch 00001: loss improved from inf to 1.29099, saving model to model-1.29.h5

409/409 [=====] - 34s 84ms/sample - loss: 1.2910 - dice\_coefficient: 0.

Epoch 2/10

408/409 [=====>.] - ETA: 0s - loss: 0.7735 - dice\_coefficient: 0.

Epoch 00002: loss improved from 1.29099 to 0.77362, saving model to model-0.77.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.7736 - dice\_coefficient: 0.

Epoch 3/10

408/409 [=====>.] - ETA: 0s - loss: 0.6429 - dice\_coefficient: 0.

Epoch 00003: loss improved from 0.77362 to 0.64259, saving model to model-0.64.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.6426 - dice\_coefficient: 0.

Epoch 4/10

408/409 [=====>.] - ETA: 0s - loss: 0.5626 - dice\_coefficient: 0.

Epoch 00004: loss improved from 0.64259 to 0.56274, saving model to model-0.56.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.5627 - dice\_coefficient: 0.

Epoch 5/10

408/409 [=====>.] - ETA: 0s - loss: 0.5050 - dice\_coefficient: 0.

Epoch 00005: loss improved from 0.56274 to 0.50588, saving model to model-0.51.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.5059 - dice\_coefficient: 0.

Epoch 6/10

408/409 [=====>.] - ETA: 0s - loss: 0.4799 - dice\_coefficient: 0.

Epoch 00006: loss improved from 0.50588 to 0.47993, saving model to model-0.48.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.4799 - dice\_coefficient: 0.

Epoch 7/10

408/409 [=====>.] - ETA: 0s - loss: 0.4601 - dice\_coefficient: 0.

Epoch 00007: loss improved from 0.47993 to 0.45990, saving model to model-0.46.h5

409/409 [=====] - 23s 57ms/sample - loss: 0.4599 - dice\_coefficient: 0.

Epoch 8/10

408/409 [=====>.] - ETA: 0s - loss: 0.4464 - dice\_coefficient: 0.

Epoch 00008: loss improved from 0.45990 to 0.44667, saving model to model-0.45.h5

409/409 [=====] - 24s 58ms/sample - loss: 0.4467 - dice\_coefficient: 0.

Epoch 9/10

408/409 [=====>.] - ETA: 0s - loss: 0.4301 - dice\_coefficient: 0.

Epoch 00009: loss improved from 0.44667 to 0.43122, saving model to model-0.43.h5

409/409 [=====] - 24s 58ms/sample - loss: 0.4312 - dice\_coefficient: 0.

Saved successfully!

408/409 [=====>.] - ETA: 0s - loss: 0.4161 - dice\_coefficient: 0.

Epoch 00010: loss improved from 0.43122 to 0.41568, saving model to model-0.42.h5

409/409 [=====] - 24s 58ms/sample - loss: 0.4157 - dice\_coefficient: 0.

<tensorflow.python.keras.callbacks.History at 0x7fcea3354400>

### ▼ Get the predicted mask for a sample image (3 marks)

```
n = 10
sample_image = X_train[n]
#### Add your code here ####
print(sample_image.shape)
sample_image_reshaped = np.reshape(sample_image, (1, sample_image.shape[0], sample_image.shape[1], sample_image.shape[2]))
print(sample_image_reshaped.shape)
predicted_mask = model.predict(sample_image_reshaped)
```

```
↳ (224, 224, 3)
   (1, 224, 224, 3)
```

```
print(predicted_mask.shape)
```

Saved successfully!

```
print(predicted_mask, masks[n].shape)
```

```
↳ (1, 224, 224)
   (224, 224)
   (224, 224)
```

### ▼ Impose the mask on the image (3 marks)

```
#### Add your code here ####
#pyplot.imshow(masks[n])
pyplot.imshow(predicted_mask_reshaped)
```

```
↳
```